

# Introduction to Operating Systems (SWE3004)

## Project #2 – CPU Scheduling



Beomseok Nam (남범석)

# Xv6 Process

- Process states (procstate in proc.h)

```
52 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
53
54 // Per-process state
55 struct proc {
56     uint sz;                // Size of process memory (bytes)
57     pde_t* pgdir;           // Page table
58     char *kstack;           // Bottom of kernel stack for this process
59     enum procstate state;    // Process state
60     int pid;                // Process ID
```

- — UNUSED: Not used
- — EMBRYO: Newly allocated (not ready for running yet)
- — SLEEPING: Waiting for I/O, child process, or time
- — RUNNABLE: Ready to run
- — RUNNING: Running on CPU
- — ZOMBIE: Exited

# Xv6 Process Scheduler

- main() in main.c

```
17 int
18 main(void)
19 {
20     userinit();    // first user process
21     // Finish setting up this processor in mpmain.
22     mpmain();
23 }
```

- mpmain() in main.c

```
54 // Common CPU setup code.
55 static void
56 mpmain(void)
57 {
58     cprintf("cpu%d: starting\n", cpu->id);
59     idtinit();    // load idt register
60     xchg(&cpu->started, 1); // tell startothers() we're up
61     scheduler();  // start running processes
62 }
```

# Xv6 Process Scheduler

- scheduler() in proc.c
  - Round-robin fashion

```
265 void
266 scheduler(void)
267 {
268     struct proc *p;
269
270     for(;;){
271         // Enable interrupts on this processor.
272         sti();
273
274         // Loop over process table looking for process to run.
275         acquire(&ptable.lock);
276         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
277             if(p->state != RUNNABLE)
278                 continue;
279
280             // Switch to chosen process. It is the process's job
281             // to release ptable.lock and then reacquire it
282             // before jumping back to us.
283             proc = p;
284             switchuvm(p);
285             p->state = RUNNING;
286             swtch(&cpu->scheduler, proc->context);
287             switchkvm();
288
289             // Process is done running for now.
290             // It should have changed its p->state before coming back.
291             proc = 0;
292         }
293         release(&ptable.lock);
294     }
295 }
296 }
```

# Xv6 Process Scheduler

- `swtch()` in `swtch.S`

```
8  .globl swtch
9  swtch:
10     movl 4(%esp), %eax
11     movl 8(%esp), %edx
12
13     # Save old callee-save registers
14     pushl %ebp
15     pushl %ebx
16     pushl %esi
17     pushl %edi
18
19     # Switch stacks
20     movl %esp, (%eax)
21     movl %edx, %esp
22
23     # Load new callee-save registers
24     popl %edi
25     popl %esi
26     popl %ebx
27     popl %ebp
28     ret
```

## Xv6 Entering Scheduler

- sched() in proc.c

```
298 // Enter scheduler. Must hold only ptable.lock
299 // and have changed proc->state.
300 void
301 sched(void)
302 {
303     int intena;
304
305     if(!holding(&ptable.lock))
306         panic("sched ptable.lock");
307     if(cpu->ncli != 1)
308         panic("sched locks");
309     if(proc->state == RUNNING)
310         panic("sched running");
311     if(readeflags() & FL_IF)
312         panic("sched interruptible");
313     intena = cpu->intena;
314     swtch(&proc->context, cpu->scheduler);
315     cpu->intena = intena;
316 }
```

# Xv6 Entering Scheduler

- When?

1. Exiting process (exit() in proc.c)

```
208 // Jump into the scheduler, never to return.
209 proc->state = ZOMBIE;
210 sched();
```

2. Sleeping process (sleep() in proc.c)

```
371 // Go to sleep.
372 proc->chan = chan;
373 proc->state = SLEEPING;
374 sched();
```

# Xv6 Entering Scheduler

## ■ When?

### 3. Yielding CPU due to timer interrupt

- trap() in trap.c

```
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(proc && proc->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER)
106     yield();
```

- yield() in proc.c

```
318 // Give up the CPU for one scheduling round.
319 void
320 yield(void)
321 {
322     acquire(&ptable.lock); //DOC: yieldlock
323     proc->state = RUNNABLE;
324     sched();
325     release(&ptable.lock);
326 }
```



## Project Assignment #2

- Implement **MLFQ scheduler** on xv6
  - The lower nice value, the higher priority
  - The highest priority process is selected for next running
  - **Tiebreak: round-robin fashion**
  - **More clarifications on the requirements will be announced later.**
  
- Entering scheduler when
  1. Exiting process
  2. Sleeping process
  3. Yielding CPU
  4. **Changing priority**

## Template Code

- git clone <https://github.com/jinsoox/xv6-skku.git> -b pa2
  
- Modifications
  - halt system call
    - Halt xv6 program
  - make tarball
    - Compress your source codes into one .tar.gz file for submission
    - You should enter your ID & project no. on Makefile
  - CPUS=1
  - Ignore to yield CPU on clock tick
  - yield system call
    - Yield CPU